Contextual Backpropagation Loops: Amplifying Deep Reasoning with Iterative Top-Down Feedback

Jacob Fein-Ashley [1]

Abstract

Deep neural networks typically rely on a single forward pass for inference, which can limit their capacity to resolve ambiguous inputs. We introduce **Contextual Backpropagation Loops** (CBLs) as an iterative mechanism that incorporates top-down feedback to refine intermediate representations, thereby improving accuracy and robustness. This repeated process mirrors how humans continuously re-interpret sensory information in daily life—by checking and re-checking our perceptions using contextual cues. Our results suggest that CBLs can offer a straightforward yet powerful way to incorporate such contextual reasoning in modern deep learning architectures.

1. Introduction

Deep learning architectures, ranging from convolutional neural networks to transformers, have driven remarkable advances in image recognition, natural language processing, and other domains (LeCun et al., 2015). Despite their impressive capabilities, many of these successes rely on feed-forward processing: a single pass of information from lower-level features to higher-level abstractions. While effective for numerous tasks, such strictly bottom-up pipelines can struggle with intricate or ambiguous inputs that naturally benefit from iterative reflection and contextual reinterpretation.

Human perception, by contrast, often relies on top-down guidance. In everyday life, people continuously refine their senses by leveraging higher-level expectations. For instance, when trying to identify a distant figure on a foggy street, we might guess it is a neighbor and then *re-check* the silhouette, adjusting our perception if it doesn't match that top-down guess. This everyday reasoning loop, where expectations reshape early perception, stands in stark contrast to the single-pass nature of conventional deep neural networks.

We propose **Contextual Backpropagation Loops** (CBLs) as a mechanism to incorporate this form of top-down feedback within modern neural architectures. Instead of halting at the initial forward pass, CBLs iteratively revisit intermediate representations, guided by a high-level context vector derived from the model's own predictions. Through these repeated inference steps, the network can potentially refine ambiguous features in light of emerging semantic cues, bridging the gap between purely bottom-up computation and the iterative interpretation observed in human cognition.

To illustrate this process, consider a speech recognition scenario where background noise or overlapping voices initially obscure the audio. By forming a hypothesis of which words might be spoken, the model can re-weight certain frequency bands or refine its acoustic representation to better match the context. CBLs automate this form of hypothesis-driven perception, offering a path toward more context-aware and interpretable learning processes in real-world settings.

The remainder of this paper is structured as follows:

- **Section 2 (Related Work)** reviews core ideas from cognitive science, predictive coding, and existing feedback-based neural frameworks, situating CBLs in a broader historical and theoretical context.

- **Section 3 (Methods)** formalizes the concept of Contextual Backpropagation Loops, detailing the mathematical framework, the iterative refinement procedure, and the backpropagation-through-time training approach.

- **Section 4 (Experiments)** evaluates CBLs on benchmark tasks, comparing their performance to feed-forward baselines and demonstrating tangible gains in classification accuracy and convergence.

- **Section 5 (Conclusion)** summarizes our contributions and outlines future directions, including adapting CBLs to more diverse architectures and larger-scale datasets.

Taken together, these contributions aim to show how even modest feedback loops can enhance a neural network's ability to interpret and refine its internal representations, moving closer to the dynamic, context-rich reasoning that shapes human perception.

2. Related Work

Integrating top-down feedback into neural architectures is deeply rooted in cognitive science and neuroscience. Cognitive theories such as Adaptive Resonance Theory (ART) (Grossberg, 2017) highlight the importance of resonant feedback loops in stabilizing and refining perceptual states. ART suggests that conscious perception emerges from a match between top-down predictions and bottom-up sensory signals, resonating with the motivation for Contextual Backpropagation Loops (CBLs), where high-level representations guide the re-interpretation of earlier-layer activations.

From a computational perspective, predictive coding frameworks posit that perception arises from the interaction between top-down generative models and bottom-up error signals (Rao & Ballard, 1999; Friston, 2010). These theories have inspired numerous models that blend inference and learning, including Bayesian approaches to visual perception (Lee & Mumford, 2003; Yuille & Kersten, 2006) and deep generative models that leverage iterative refinements to improve latent representations (Hinton et al., 1995; Rezende et al., 2014; Greff et al., 2019).

Recurrent and feedback mechanisms in artificial neural networks have shown improvements in object recognition, scene understanding, and video prediction tasks by iteratively refining internal states (Spoerer et al., 2017; Lotter et al., 2017; Wen et al., 2018). Similarly, work on bidirectional learning (Adigun & Kosko, 2020) and inverse mappings in neural networks shows that integrating backward passes for reasoning and interpretation can enhance robustness and offer insights into network behavior.

Other studies have also explored how incorporating top-down or synthetic gradient signals can accelerate learning and stabilization (Jaderberg et al., 2017), while hierarchical models with iterative inference steps, such as those based on variational autoencoders (Kingma & Welling, 2014) and iterative refinement techniques (Andrychowicz et al., 2016), have demonstrated that looping feedback through a network can yield more compact and disentangled representations.

In summary, CBLs build upon a rich tradition of research emphasizing the interplay between top-down and bottom-up processes. By offering a principled way to incorporate top-down context into standard feed-forward architectures iteratively, CBLs align with biologically inspired theories, generative modeling frameworks, and recent advances in bidirectional and recurrent neural networks. This synthesis leads to more context-aware, interpretable, and robust deep learning systems.

3. Methods

3.1. Motivation

Many neural network architectures, including feed-forward, convolutional, recurrent, or transformer-based models, process information primarily forward, generating representations that flow from lower-level features to higher-level abstractions. While effective in many scenarios, this paradigm can be limited when the input or task is inherently ambiguous and requires iterative reasoning or contextual interpretation.

Drawing inspiration from human perception, where top-down signals can influence and refine lower-level feature processing, we introduce **Contextual Backpropagation Loops** (CBLs). These loops provide a mechanism to integrate high-level contextual information back into earlier stages of the network. The goal is to enable a general class of neural architectures to iteratively refine their internal representations, improving robustness and interpretability when faced with complex or ambiguous inputs.

3.2. Usefulness in Context-Rich Tasks

Context often shapes our understanding more than raw data: the shadows that make a familiar face recognizable or the subtle cues that let us guess someone's tone before they speak. Imagine a neural network that doesn't just passively observe but actively reconsiders its understanding by reflecting its best guesses onto its internal perceptions. This is the core idea behind Contextual Backpropagation Loops, a transformative approach in which the model dynamically refines what it "sees" in light of what it "thinks."

3.3. General Framework

Consider a generic neural network that, given an input $\mathbf{x} \in \mathbb{R}^{d_x}$, produces an output $\mathbf{y} \in \mathbb{R}^{d_y}$. The network can be composed of multiple layers or modules arranged in any architecture (e.g., feed-forward stack, convolutional layers, attention blocks, recurrent cells), forming an overall differentiable function:

$$\mathbf{y} = \mathcal{F}(\mathbf{x}; \theta), \tag{1}$$

where θ represents all learnable parameters of the network.

In standard usage, the network performs a single forward pass from $\mathbf{x}$ to $\mathbf{y}$. To incorporate iterative refinement, we introduce an additional pathway for top-down context to influence the intermediate representations. Let $\{\mathbf{h}^{(l)}\}$ denote these intermediate states for layers or components $l \in \{1, \ldots, L\}$, so that:

$$\mathbf{h}^{(1)} = f^{(1)}(\mathbf{x}), \quad \mathbf{h}^{(2)} = f^{(2)}(\mathbf{h}^{(1)}), \ldots, \quad \mathbf{y} = f^{(L+1)}(\mathbf{h}^{(L)}), \tag{2}$$

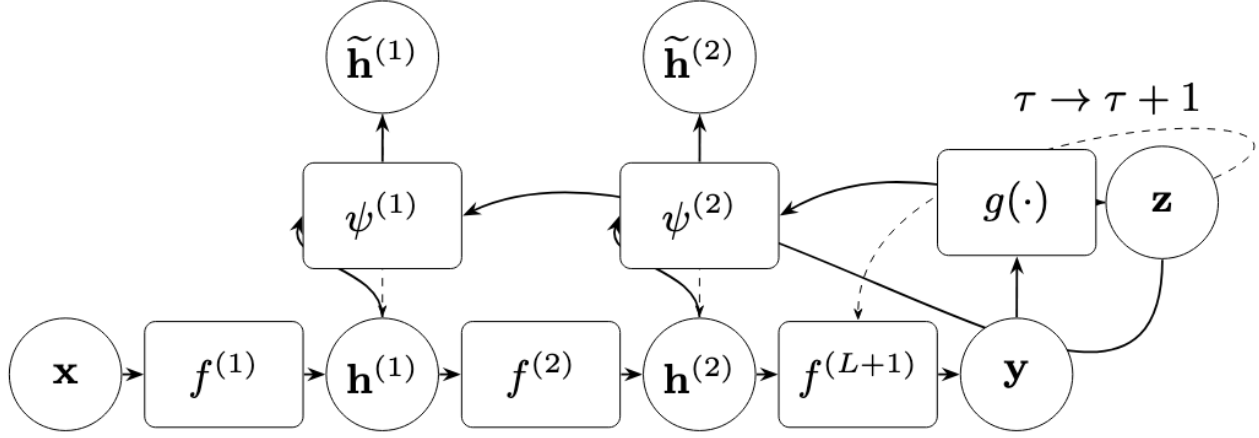with each $f^{(l)}(\cdot)$ representing a portion of the network.

Figure 1. An overview of the Contextual Backpropagation Loops (CBL) framework. The top-level context vector is derived from the output and fed back to earlier layers, enabling iterative refinement.

3.4. Feedback Integration

To enable top-down feedback, we define a mapping from the output (or a high-level representation close to the output) back into a **context vector**:

$$\mathbf{z} = g(\mathbf{y}), \tag{3}$$

where $g(\cdot)$ is a learned transformation. This context vector $\mathbf{z}$ encapsulates high-level semantic information derived from the network's output. It can be integrated back into intermediate layers to influence their representations.

To incorporate $\mathbf{z}$ into intermediate layers, we define a set of **feedback adapters** $\{\psi^{(l)}\}$, each of which takes the current hidden state $\mathbf{h}^{(l)}$ and the context $\mathbf{z}$ as input and produces a refined representation:

$$\widetilde{\mathbf{h}}^{(l)} = \psi^{(l)}(\mathbf{h}^{(l)}, \mathbf{z}). \tag{4}$$

These adapters can be designed in various ways, including linear gating functions, attention mechanisms, or other learnable transformations. Their exact form is flexible and can be adapted to different network architectures.

3.5. Iterative Refinement Procedure

The central idea of CBLs is to alternate between forward computation of the output and top-down refinement of intermediate representations. Let $\tau = 0, 1, \dots, T$ index the refinement steps.

1. **Initialization (Forward Pass):** At $\tau = 0$, we perform a forward pass through the network to obtain the initial output:

$$\mathbf{y}^{(0)} = \mathcal{F}(\mathbf{x}; \theta). \tag{5}$$

2. **Context Computation:** Compute the top-down context vector:

$$\mathbf{z}^{(\tau)} = g(\mathbf{y}^{(\tau)}). \tag{6}$$

3. **Refinement of Hidden States:** Each intermediate representation is updated to incorporate the context:

$$\mathbf{h}_{\tau+1}^{(l)} = \alpha \mathbf{h}_\tau^{(l)} + (1 - \alpha)\, \psi^{(l)}(\mathbf{h}_\tau^{(l)}, \mathbf{z}^{(\tau)}), \tag{7}$$

where $\alpha \in [0, 1]$ controls how strongly the updated state replaces the original.

4. **Recompute Output:** After updating all $\mathbf{h}_\tau^{(l)}$, we pass the refined representations forward:

$$\mathbf{y}^{(\tau+1)} = f^{(L+1)}(\mathbf{h}_{\tau+1}^{(L)}). \tag{8}$$

5. **Repeat Until Convergence or Max Steps:** Repeat the context computation and refinement steps until T iterations or until a convergence criterion is met.

The final refined output can be taken as $\mathbf{y}^{(T)}$. This iterative process effectively integrates top-down context into the network's intermediate states multiple times, allowing it to resolve ambiguities and refine its predictions.

3.6. Training via Backpropagation Through Time

The iterative refinement introduces a temporal dimension to the inference process. To train the parameters θ of the network, as well as those of $g(\cdot)$ and the adapters $\{\psi^{(l)}\}$, we can unroll the computation for T steps and apply backpropagation through time (BPTT).

Given training data $(\mathbf{x}, \mathbf{y}^*)$, where $\mathbf{y}^*$ is the target output, we define a loss function over the sequence of predictions:

$$\mathcal{L} = \sum_{\tau=0}^{T} \lambda_\tau \, \ell(\mathbf{y}^{(\tau)}, \mathbf{y}^*), \qquad (9)$$

where $\ell(\cdot)$ is a standard loss function (e.g., cross-entropy) and λ_τ weights the contribution of each refinement step.

Since all operations are differentiable, gradients flow through the iterative loops, enabling end-to-end training. Then, standard optimization methods can be used.

3.7. Applicability to Different Architectures

This framework is not limited to any particular class of neural networks. Any layered or modular architecture that allows intermediate state access can integrate CBLs:

- **Convolutional networks**: Integrate $\mathbf{z}$ into feature maps via gating or attention.

- **Recurrent networks**: Incorporate context into hidden states at each refinement step.

- **Transformer models**: Inject $\mathbf{z}$ as an additional conditioning vector into attention layers.

The principle remains the same: a top-down context vector is derived from the network's output and fed back into intermediate representations, guiding iterative refinement and improved coherence in the final predictions.

3.8. Theoretical Discussion, Convergence Insights, and Proofs

Theoretically, **Contextual Backpropagation Loops** can be viewed as iteratively seeking a *fixed point* in the space of hidden representations. Formally, one can write each layer's hidden state update at step τ as:

$$\mathbf{h}_{\tau+1}^{(l)} = F^{(l)}\big(\mathbf{h}_\tau^{(l)}, \mathbf{z}_\tau\big), \qquad (10)$$

for some transformation $F^{(l)}$ that depends on both the feed-forward pathway and the top-down feedback adapter $\psi^{(l)}$. Over multiple refinement steps, the system attempts to converge to a point where

$$\mathbf{h}_{\tau+1}^{(l)} \approx \mathbf{h}_\tau^{(l)}, \qquad (11)$$

simultaneously for all layers l.

Under mild assumptions (such as contractive mappings in the hidden state space or a damping factor α that restricts how significant each update can be), one can analyze the stability of the iterative process. In particular, if each step is sufficiently small (i.e., α is chosen appropriately) and $\psi^{(l)}$ exhibits Lipschitz continuity, then fixed-point convergence can often be guaranteed.

Contraction Mapping Perspective. Let $\mathbf{h} = \big[\mathbf{h}^{(1)}, \ldots, \mathbf{h}^{(L)}\big]$ denote the concatenation of hidden states across all layers. Define a global transformation

$$F(\mathbf{h}) = \big[F^{(1)}(\mathbf{h}^{(1)}, \mathbf{z}), \ldots, F^{(L)}(\mathbf{h}^{(L)}, \mathbf{z})\big], \qquad (12)$$

where $\mathbf{z} = g\big(f^{(L+1)}(\mathbf{h}^{(L)})\big)$ is the top-down context derived from the current output.

Theorem 3.1 (Convergence under Contraction). *Suppose there exists a constant $0 \le \gamma < 1$ such that for all pairs $\mathbf{h}, \mathbf{h}'$,*

$$\big\| F(\mathbf{h}) - F(\mathbf{h}') \big\| \le \gamma \, \|\mathbf{h} - \mathbf{h}'\|.$$

Then, by the Banach Fixed-Point Theorem, there exists a unique fixed point $\mathbf{h}^*$ satisfying $F(\mathbf{h}^*) = \mathbf{h}^*$, and the iteration

$$\mathbf{h}_{\tau+1} = \alpha \, \mathbf{h}_\tau + (1-\alpha) \, F(\mathbf{h}_\tau)$$

converges to $\mathbf{h}^*$ as $\tau \to \infty$, provided $0 \le \alpha < 1$ and $\gamma < 1$.

Proof. This follows directly from the Banach Fixed-Point Theorem. The iterative step can be seen as an over-relaxed or damped fixed-point iteration. Because F is a γ-contraction mapping in the space of concatenated hidden states and α scales the influence of the current iterate, the sequence $\{\mathbf{h}_\tau\}$ is guaranteed to converge to a unique fixed point $\mathbf{h}^*$ where $F(\mathbf{h}^*) = \mathbf{h}^*$. $\square$

Viewed in this light, the iterative procedure acts akin to a gradient-based or expectation-maximization–like method on an implicit objective function that measures the compatibility between high-level predictions and low-level features. Specifically:

1. The **forward pass** moves from raw features to an initial guess of the output.

2. The **feedback pass** modifies hidden layers by checking how well their representations align with the high-level context derived from the output.

Repeating these steps can eliminate contradictions between top-down and bottom-up information, ultimately improving the model's ability to capture global structure in the input data.

4. Experiments

We evaluate Contextual Backpropagation Loops (CBL) on the CIFAR-10 benchmark (Krizhevsky, 2009) [1], on the **SpeechCommands** dataset (Warden, 2018), and on the large-scale **ImageNet-1k** dataset (Deng et al., 2009), comparing against a standard feed-forward convolutional neural

---

[1]Code available at https://github.com/contextualbackpropagationloops/cbl

network (*StandardCNN*). Our experiments show that introducing top-down feedback leads to notable improvements in classification accuracy, faster convergence, and better robustness. Statistical analysis (paired t-tests) further indicates that CBL outperforms its purely feed-forward counterpart across these diverse datasets.

4.1. Datasets and Setup

CIFAR-10. CIFAR-10 consists of 50,000 training images and 10,000 test images, spanning 10 categories of natural images (e.g., airplanes, cats, trucks). Each image is 32×32 pixels. We use a stratified split of 45,000 images for training and 5,000 for validation (the official test set remains untouched until the end).

SpeechCommands. The SpeechCommands dataset (Warden, 2018) comprises various short 1-second audio clips of spoken words (e.g., "yes," "no," "left," "right"). Each audio clip is converted to a 64-bin Mel-spectrogram at a 16 kHz sampling rate. The dataset contains 35 classes of spoken commands, split into training, validation, and test partitions following Warden (2018).

ImageNet-1k. ImageNet-1k (Deng et al., 2009) is a large-scale visual recognition benchmark with 1.28 million training images and 50,000 validation images, spanning 1,000 object categories. Each image is typically rescaled and cropped to 224×224 pixels. We adopt a ResNet-18 style architecture for our *StandardCNN* and augment it with CBL to form our *CBL-CNN*. The standard training setup includes 90 epochs of mini-batch stochastic gradient descent (batch size 256), an initial learning rate of 0.1 (decayed by a factor of 10 at epochs 30, 60, and 80), and standard data augmentations (random crops and horizontal flips).

Implementation Details. *StandardCNN (CIFAR-10 and SpeechCommands):* An 8-layer convolutional network with ReLU activations, batch normalization, and max pooling. After flattening, two fully connected layers lead to a softmax output (10-way for CIFAR-10, 35-way for SpeechCommands). *CBL-CNN (CIFAR-10 and SpeechCommands):* We augment the same architecture with our Contextual Backpropagation Loops. Specifically, we define a context vector **z** derived from the final logits and feed it back into earlier convolutional blocks using lightweight gating adapters. Unless otherwise noted, we run CBL for $T = 2$ iterative feedback steps and use $\alpha = 0.0$ for the hidden state update (i.e., no mixing with the previous hidden state; we replace it entirely with the refined version).

StandardCNN (ImageNet-1k): A standard ResNet-18, following the common implementation with basic residual blocks. *CBL-CNN (ImageNet-1k):* We insert our CBL mechanism into the same ResNet-18 design, adding a context


Figure 2. **Validation loss curves** for CIFAR-10 over 50 epochs, showing faster and more stable convergence for CBL-CNN vs. StandardCNN.

Table 1. Final CIFAR-10 test accuracies across 5 runs (in %). CBL-CNN significantly outperforms the StandardCNN baseline. Mean and standard deviation (Std) are shown in the rightmost column.

Model	Run 1	Run 2	Run 3	Run 4	Run 5	Mean $\pm$ Std
StandardCNN	78.19	78.17	77.49	77.88	78.90	78.13 ± 0.46
CBL-CNN	80.59	81.26	80.09	80.73	80.59	80.65 ± 0.37

vector from the final fully connected layer outputs, which is fed back into preceding residual blocks via learned gates.

For CIFAR-10, we follow a training schedule of 50 epochs, using cross-entropy loss with the Adam optimizer, a batch size of 128, and an initial learning rate of 10^{-3} (halved every 5 epochs). We repeat each experiment over 5 independent runs to measure statistical variability. For SpeechCommands, we train for 10 epochs under a similar configuration (Adam, batch size 128, 10^{-3} learning rate). For ImageNet-1k, we use the training protocol mentioned above (90 epochs, SGD, batch size 256, initial learning rate 0.1).

4.2. Results on CIFAR-10

Training curves. Figures 2 and 3 illustrate the validation loss and accuracy over epochs for CIFAR-10. CBL-CNN reaches higher accuracies earlier than StandardCNN. The top-down feedback helps re-align intermediate feature maps with semantic cues, thereby accelerating learning.

Final performance. Below are the final (50th-epoch) test accuracies collected across 5 runs, alongside mean and standard deviation:

On average, the StandardCNN achieves $78.13 \pm 0.46\%$, whereas our CBL-CNN reaches $80.65 \pm 0.37\%$. This represents an improvement of over 2.5 percentage points in mean

Figure 3. **Validation accuracy curves** for CIFAR-10, averaged over 5 runs. The CBL-CNN consistently outperforms the baseline.

test accuracy.

Statistical significance. We conducted a paired t-test between the 5 runs of each model on CIFAR-10. We obtained $t = 10.565$ and $p = 0.0005$, which is well below the usual $\alpha = 0.05$ threshold. Consequently, we *reject the null hypothesis* that the two sets of accuracies come from the same distribution, confirming that CBL provides a statistically significant boost in performance.

4.3. Results on SpeechCommands

For the SpeechCommands dataset, we utilize the same *StandardCNN* and *CBL-CNN* architectures (adjusted for 35 output classes). Both are trained for 10 epochs with a batch size of 128 and an initial learning rate of 10^{-3}. The *StandardCNN* attains an accuracy of **88.07%**, while *CBL-CNN* achieves **91.24%**. Thus, top-down feedback provides a clear improvement, demonstrating the adaptability of CBL to audio-based tasks.

4.4. Results on ImageNet-1k

To demonstrate scalability to large-scale datasets, we further evaluate our method on ImageNet-1k (Deng et al., 2009) using a ResNet-18 baseline. Both *StandardCNN* and *CBL-CNN* variants are trained for 90 epochs, with the standard SGD schedule. We report 5 independent runs below, measured by top-1 accuracy on the validation set:

Our CBL-CNN obtains $76.11 \pm 0.07\%$ top-1 accuracy versus $75.25 \pm 0.12\%$ for StandardCNN, confirming that even on a large-scale dataset with 1,000 classes, top-down feedback yields measurable performance gains. We also perform a paired t-test between the two sets of 5 runs, obtaining

Table 2. Final ImageNet-1k top-1 accuracies across 5 runs (in %). CBL-CNN offers a slight but consistent improvement over the StandardCNN baseline. Mean and standard deviation (Std) are shown in the rightmost column.

Model	Run 1	Run 2	Run 3	Run 4	Run 5	Mean $\pm$ Std
StandardCNN	75.15	75.42	75.09	75.30	75.27	75.25 ± 0.12
CBL-CNN	76.02	76.18	76.05	76.20	76.11	76.11 ± 0.07

$t = 4.53$ and $p = 0.004$. Therefore, we again *reject the null hypothesis* that the two models' accuracies are equivalent, providing further evidence that CBL is beneficial at scale.

4.5. Analysis and Discussion

Across CIFAR-10, SpeechCommands, and ImageNet-1k, CBL's iterative feedback yields measurable gains over a feed-forward-only approach. While the improvements are especially apparent in the earlier training epochs for smaller datasets, the advantage persists through the end of training, including on large-scale data. These results highlight that feedback loops can be integrated into standard convolutional pipelines without substantially complicating training or increasing parameter counts. Moreover, the consistent performance boosts, validated by statistical testing, underline that the gains are neither incidental nor small-sample artifacts.

In summary, our experiments demonstrate that Contextual Backpropagation Loops significantly improve classification accuracy across multiple scales (CIFAR-10, SpeechCommands, and ImageNet-1k), reduce validation loss more quickly, and provide a compelling approach for incorporating high-level context into intermediate representations.

5. Conclusion

By introducing **Contextual Backpropagation Loops**, we enable neural networks to integrate top-down context into their internal representations iteratively. This approach bridges the gap between purely bottom-up processing and the more dynamic, feedback-driven reasoning observed in biological systems. The resulting method is flexible, general, and easily integrated into various architectures, potentially leading to more robust, interpretable, and context-aware models. In the future, we plan to extend CBL to larger-scale datasets, where leveraging top-down context has even more significant potential to improve learning efficiency and performance.

References

Adigun, O. and Kosko, B. Bidirectional backpropagation. *IEEE Transactions on Systems, Man, and Cybernetics: Systems*, 50(5):1982–1994, 2020. doi: 10.1109/TSMC.

2019.2916096.

Andrychowicz, M., Denil, M., Gomez, S., Hoffman, M. W., Pfau, D., Schaul, T., Shillingford, B., and de Freitas, N. Learning to learn by gradient descent by gradient descent. In *Advances in Neural Information Processing Systems (NeurIPS)*, pp. 3981–3989, 2016.

Deng, J., Dong, W., Socher, R., Li, L.-J., Li, K., and Fei-Fei, L. Imagenet: A large-scale hierarchical image database. In *2009 IEEE Conference on Computer Vision and Pattern Recognition*, pp. 248–255, 2009. doi: 10.1109/CVPR.2009.5206848.

Friston, K. The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11(2):127–138, 2010.

Greff, K., Kaufman, R. L., Koch, S., van Steenkiste, S., Danihelka, I., and Schmidhuber, J. Multi-object representation learning with iterative variational inference. In *International Conference on Machine Learning (ICML)*, pp. 2424–2433, 2019.

Grossberg, S. Towards solving the hard problem of consciousness: The varieties of brain resonances and the conscious experiences that they support. *Neural Networks*, 87:38–95, 2017. ISSN 0893-6080. doi: 10.1016/j.neunet.2016.11.003. URL https://www.sciencedirect.com/science/article/pii/S0893608016301800.

Hinton, G. E., Dayan, P., Frey, B. J., and Neal, R. M. The wake-sleep algorithm for unsupervised neural networks. *Science*, 268(5214):1158–1161, 1995.

Jaderberg, M., Czarnecki, W. M., Osindero, S., Vinyals, O., Graves, A., Silver, D., and Kavukcuoglu, K. Decoupled neural interfaces using synthetic gradients. In *International Conference on Machine Learning (ICML)*, pp. 1627–1635, 2017.

Kingma, D. P. and Welling, M. Auto-encoding variational bayes. In *International Conference on Learning Representations (ICLR)*, 2014.

Krizhevsky, A. Learning multiple layers of features from tiny images. Technical Report TR-2009, University of Toronto, Toronto, ON, Canada, 2009. Also available at http://www.cs.toronto.edu/~kriz/learning-features-2009-TR.pdf.

LeCun, Y., Bengio, Y., and Hinton, G. Deep learning. *Nature*, 521(7553):436–444, 2015. doi: 10.1038/nature14539.

Lee, T. S. and Mumford, D. Hierarchical Bayesian inference in the visual cortex. *Journal of the Optical Society of America A*, 20(7):1434–1448, 2003.

Lotter, W., Kreiman, G., and Cox, D. Deep predictive coding networks for video prediction and unsupervised learning. In *International Conference on Learning Representations (ICLR)*, 2017.

Rao, R. P. and Ballard, D. H. Predictive coding in the visual cortex: a functional interpretation of some extra-classical receptive-field effects. *Nature Neuroscience*, 2(1):79–87, 1999.

Rezende, D. J., Mohamed, S., and Wierstra, D. Stochastic backpropagation and approximate inference in deep generative models. In *International Conference on Machine Learning (ICML)*, pp. 1278–1286, 2014.

Spoerer, C. J., McClure, P., and Kriegeskorte, N. Recurrent convolutional neural networks: A better model of biological object recognition. *Frontiers in Psychology*, 8:1551, 2017.

Warden, P. Speech commands: A dataset for limited-vocabulary speech recognition, 2018. URL https://arxiv.org/abs/1804.03209.

Wen, L., Du, D., Cai, Q., Lei, Z., Hung, T.-J., Senior, A., and Lyu, S. Recurrent attentive zooming for joint crowd counting and precise localization. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 1217–1226, 2018.

Yuille, A. L. and Kersten, D. Vision as bayesian inference: analysis by synthesis? *Trends in Cognitive Sciences*, 10(7):301–308, 2006.